

# Kubernetes

## 0. Terminology

- container runtime
- kubelet

## 1. Docker Basics

Steps - create docker image - docker build -t quickstart-image . - docker tag quickstart-image gcr.io/[PROJECT-ID]/quickstart-image:tag1 - docker push gcr.io/[PROJECT-ID]/quickstart-image:tag1 - docker pull gcr.io/[PROJECT-ID]/quickstart-image:tag1 - gcloud container images delete gcr.io/[PROJECT-ID]/quickstart-image:tag1 --force-delete-tags - kubernetes yaml file

```
# docker file syntax

FROM [--platform=<platform>] <image>[:<tag>] [AS <name>]

ENTRYPOINT: set image's main command
CMD
ENV: provide environment variables
ARG
RUN
```

how to speed up image build? troubleshooting

```
# docker-compose cannot kill network
docker network inspect <network_name>
docker network disconnect -f <network_name> <container_name>
```

## 2. Core Concepts

Goals: - Understand Kubernetes API primitives - Create and configure basic Pods

Components - api server - etcd - scheduler - kubelet - controller - container runtime

### 2.1 Workloads

#### 2.1.1 Pod

the smallest unit, can have more than one containers running

```
# kubectl cmd

kubectl run <pod_name> --image <image_name>

kubectl create -f example.yaml

kubectl get pods -n <namespace>
kubectl get pod <pod-name> -o yaml > pod-definition.yaml # get pod definition file
kubectl edit pod <pod-name>
kubectl get replicationcontroller

# use selector to retrieve info from matched pods

kubectl get pods -l environment=production,tier=frontend # equality-based
kubectl get pods -l 'environment in (production),tier in (frontend)' # set-based
kubectl get pod --sort-by=.metadata.creationTimestamp -n dbt-tactical-ds -o name
```

#### 2.1.2 ReplicaSet

```
...
kind: ReplicationController
metadata:
  name:
  label:
    l_name: l_value
spec:
  template:
    <pod definition>
  replicas: 2

# replicationSet
kind: ReplicaSet
selector:
```

```
# how to scale a replicaSet?

kubectl replace -f .yaml
kubectl scale --replicas=6 -f .yaml
kubectl config set-context $(kubectl config current-context) --namespace=dev
```

### 2.1.3 Deployment

```
kind: Deployment
```

### 2.1.4 namespace

```
kind: namespaces
```

### 2.1.5 Service

```
kind: service
```

### 2.1.6 Imperative command

```
kubectl run --generator=run-pod/v1 nginx-pod --image=nginx:alpine

# expose servian from a pod
kubectl expose pod redis --port=6379 --name redis-service

# create deployment and scale
kubectl create deployment webapp --image=kodekloud/webapp-color
kubectl scale deployment/webapp --replicas=3

# flag
-o output
-l label
```

## 2.2 Configuration

Goals: - Understand ConfigMaps - Understand SecurityContexts - Define an application's resources requirements - Create & consume Secrets - Understand ServiceAccounts

patterns: - define command and arguments for a container - define environment variable - expose pod information via env\_var - distribute credential using secrets - inject information using PodPreset

```
command --> ENTRYPOINT
```

args --> CMD

```
spec:
  container:
    image: nginx
    args: ['run', 'a']
```

## 2.2.1 configMaps

---

```
kubectl get configmaps
kubectl create configmap <cm-name> --from-literal=<key>=<name>
```

# how to define configmap in container block

```
envFrom:
- configMapRef:
  name: special-config
```


## 2.2.2 Secrets, SecurityContexts

---

```
kubectl get secrets
```

# create new secret

```
kubectl create secret generic db-secret --from-literal=DB_Host=sql01 --from-literal=DB_User=root --from-literal=DB_Password=password
```



secret type - service account - secret - Opaque

# use secret in image block

```
envFrom:
- secretRef:
  name: special-config
```

# security contexts

```
kubectl exec ubuntu-sleeper whoami
```

## 2.2.3 Resource Limits

---

[Assign CPU resources ot Containers and Pods](#)

```
# container resources and limits
```

```
apiVersion: v1
kind: Pod
metadata:
  name: cpu-demo-2
  namespace: cpu-example
spec:
  containers:
  - name: cpu-demo-ctr-2
    image: vish/stress
    resources:
      limits:
        cpu: "100"
      requests:
        cpu: "100"
    args:
    - -cpus
    - "2"
```

## 2.2.4 Service Account

### Usage:

To communicate with the API server, a Pod uses a ServiceAccount containing an authentication token.

```
# service account
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: build-robot
```

```
# create via cli
```

```
kubectl apply -f - <<EOF
apiVersion: v1
kind: ServiceAccount
metadata:
  name: build-robot
EOF
```

```
# retrieve token
```

```
kubectl describe secret <dashboard-sa-secret-name>
```

```
# example
```

```
$ kubectl get secret default-token-dffkj -o yaml
```

```
apiVersion: v1
data:
  ca.crt: LS0tLS1CRU...0tLS0tCg==
  namespace: ZGVmYXVsdA==
  token: ZXlKaGJHY2...RGMU1IX2c=
kind: Secret
metadata:
  name: default-token-dffkj
  namespace: default
...
type: kubernetes.io/service-account-token
```

```

# binding roles to sa

# sa
apiVersion: v1
kind: ServiceAccount
metadata:
  name: demo-sa

# role
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: Role
metadata:
  name: list-pods
  namespace: default
rules:
  - apiGroups:
    - ''
  resources:
    - pods
  verbs:
    - list

# binding
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: RoleBinding
metadata:
  name: list-pods_demo-sa
  namespace: default
roleRef:
  kind: Role
  name: list-pods
  apiGroup: rbac.authorization.k8s.io
subjects:
  - kind: ServiceAccount
    name: demo-sa
    namespace: default

```

`ca.crt` is the Base64 encoding of the cluster certificate.

`token` is the Base64 encoding of the JWT used to authenticate against the API server.

what is the relationship between GSA and KSA?

## 2.3 multi-container pods

Goals: - Understand Multi-Container Pod design patterns

### 2.3.1 taint, toleration

taint and toleration are mostly used together, they are mainly used to manage resources for specific usage cases.

what is effect in tolerations? e.g. `NoSchedule`, `NoExecute`

```

# taint

kubectl taint nodes node01 spray=mortein:NoSchedule

```

### 2.3.2 Affinity

[Node affinity](#)

node affinity allows you to constrain which nodes your pod is eligible to be scheduled on, based on labels on the node.

### 2.3.3 Common patterns

Patterns: - sidecar - adapter - ambassador

sidecar ([logging services](#))

[Using a sidecar container with the logging agent](#)

```
# create sidecar container for storing logging

apiVersion: v1
kind: Pod
metadata: app
  name: app
  namespace: elastic-stack
  labels:
    name: app
spec:
  containers:
    - name: app
      image: kodekloud/event-simulator
      volumeMounts:
        - mountPath: /log
          name: log-volume

    - name: sidecar
      image: kodekloud/filebeat-configured
      volumeMounts:
        - mountPath: /var/log/event-simulator/
          name: log-volume

  volumes:
    - name: log-volume
      hostPath:
        path: /var/log/webapp
        type: DirectoryOrCreate
```

## 2.4 Observability

Goals: - Understand LivenessProbes and ReadinessProbes - Understand container logging - Understand how to monitor applications - Understand debugging in kubernetes

### 2.4.1 Pod Lifecycle

[Pod Lifecycle](#)

Container probes - livenessProbe - readinessProbe - startupProbe

basically it detects if the application is ready for receiving traffic

- pending
- containerCreating
- running
- Error
- completed

### 2.4.1 Logging

logging architecture

- node level logging
- cluster level logging

You should remember that native kubernetes does not support extensive logging mechanism. The managed service like GKE, EKS makes life easier at some cost.

logging agent options: - stackdriver monitoring - elastic search

```
kubectl logs -f <pod-name> <container-name>
```

### 2.4.2 Monitoring

In GKE, cloud monitoring makes the monitoring an ease to use. Apart from that, datadog is a go-to option in the industry.

## 2.5 Pod Design

Goals: - Understand Deployments and how to perform rolling updates, rollbacks - Understand Jobs and CronJobs - Understand how to use Labels, Selectors and Annotations

### 2.5.1 labels, selectors

```
# use label to retrieve info
kubectl get pods -l <label name>

# Identify the POD which is 'prod', part of 'finance' BU and is a 'frontend' tier?
# set based syntax
kubectl get pods -l 'env in (prod), bu in(finance), tier in (frontend)'
# equality based syntax
kubectl get pods -l env=prod, bu=finance, tier=frontend

# select resources based on resource fields
kubectl get pods --field-selector status.phase=Running
```

```
# metadata block has pre-defined key, including name, labels, while within labels, fields can be defined freely.

apiVersion: apps/v1
kind: Deployment
metadata:
  name: "{{ APP_NAME }}-{{ENVIRONMENT}}"
  labels:
    app: "{{ APP_NAME }}-{{ENVIRONMENT}}"
    env: {{ ENVIRONMENT }}
    app.kubernetes.io/name: "{{ APP_NAME }}-{{ENVIRONMENT}}"
    app.kubernetes.io/version: "{{ APP_VERSION }}"
    app.kubernetes.io/component: dataflow
    app.kubernetes.io/managed-by: chappie
# selector's label name should match labels field inside template
spec:
  selector:
    matchLabels:
      app: {{ APP_NAME }}-{{ENVIRONMENT}}
  template:
    metadata:
      labels:
        app: {{ APP_NAME }}-{{ENVIRONMENT}} # should match selector label
```

### 2.5.2 Rolling updates

strategy type: - RollingUpdate - 25% max unavailable, 25\$ max surge - Recreate - RollingBacks

```
kubectl edit deployment frontend
kubectl set image <deployment-name> <image-name>
kubectl rollout status <deployment-name>
kubectl rollout history <deployment-name>
kubectl rollout undo <deployment-name>

# interacting with k8s pods/node
kubectl exec --namespace=<ns> curl -- sh -c '<doing something>'
```

### 2.5.3 Job

job is basically a task to achieve certian goals. You can set `backoffLimit` which tells it to retry up to how many times until the goal is achieved.

what kind of tasks are suitable to run as job: - non-parallel job - parallel jobs with a work queue - parallel jobs with a fixed completion count

jobs: one-off run

batch jobs

restartPolicy: Never/Always

parallelism: 3

## 2.5.4 CronJob

looks like job and cronjob support different fields in spec

```
# cronjob configuration

apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: hello
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: <container_name>
              restartPolicy: OnFailure
```

e.g. `30 19 */1 * *` schedule the job to run everyday at 19:30 UTC

[cron expression validator](#)

Field name	Mandatory?	Allowed values	Allowed special characters
Seconds	Yes	0-59	* / , -
Minutes	Yes	0-59	* / , -
Hours	Yes	0-23	* / , -
Day of month	Yes	1-31	* / , - ?
Month	Yes	1-12 or JAN-DEC	* / , -
Day of week	Yes	0-6 or SUN-SAT	* / , - ?

## 2.6 Service & Networking

Goals: - Understand Services - Basic understanding of network policies

Difficulties - Define ingress/egress rules, protocol, IP, port

Services basically give pod a static address so that another pods can work upon on these backend pods.

### 2.6.1 Services

Multi-Port Services

Discovering Services - Env variables - DNS (preferred)

Publishing Services - ClusterIP - NodePort - LoadBalancer - ExternalName

Supported protocols - TCP - UDP - HTTP - PROXY protocol - SCTP



```
# sample yaml

apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  selector:
    app: MyApp # should be the name of target deployment
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

```
# expose deployments as service in a namespace
# question: how to specify the app name to be exposed via selector?
```

```
kubectl expose deployment -n ingress-space ingress-controller --type=NodePort --port=80 --name=ingress --dry-run -o yaml >ingress.yaml
```

## 2.6.2 Network Policies

ingress type: - Single Service Ingress - Simple fanout - Name based virtual hosting - TLS - Load balancing

egress

How to access deployments in other namespaces

[Nginx Ingress rewrite](#)

```
# create a new service

apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: test-ingress
  namespace: critical-space
  annotations:
    # why we need annotations here?
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
    - http:
        paths:
          - path: /pay
            backend:
              serviceName: pay-service
              servicePort: 8282
```

```
# > Create a network policy to allow traffic from the 'Internal' application only to the 'payroll-service' and 'db-service'
```

```
apiVersion: networking.k8s.io/v1
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: <policy-name>
```

```
spec:
```

```
  podSelector:
```

```
    matchLabel:
```

```
      name: internal
```

```
  policyTypes:
```

```
  - Egress
```

```
  - Ingress
```

```
  ingress:
```

```
  - {}
```

```
  egress:
```

```
  - to:
```

```
    - podSelector:
```

```
      matchLabels:
```

```
        name: <pod_name_1>
```

```
    ports:
```

```
    - protocol: TCP
```

```
      port: <port>
```

```
  - to:
```

```
    - podSelector:
```

```
      matchLabels:
```

```
        name: <pod_name_2>
```

```
    ports:
```

```
    - protocol: TCP
```

```
      port: <port>
```

```
# example ingress resources
```

```
master $ kubectl describe ingresses.networking.k8s.io ingress-wear-watch -n app-space
```

```
Name:          ingress-wear-watch
```

```
Namespace:     app-space
```

```
Address:
```

```
Default backend: default-http-backend:80 (<none>)
```

```
Rules:
```

```
Host Path Backends
```

```
----
```

```
*
```

```
    /wear    wear-service:8080 (10.32.0.2:8080)
```

```
    /watch   video-service:8080 (10.32.0.3:8080)
```

```
Annotations:
```

```
  nginx.ingress.kubernetes.io/rewrite-target: /
```

```
  nginx.ingress.kubernetes.io/ssl-redirect:    false
```

```
Events:
```

```
Type Reason Age From Message
```

```
----
```

```
Normal CREATE 6m21s nginx-ingress-controller Ingress app-space/ingress-wear-watch
```

```
Normal UPDATE 6m20s nginx-ingress-controller Ingress app-space/ingress-wear-watch
```

## 2.7 State Persistence

Goal: Understand PersistentVolumeClaims for storage

[doc link](#)

Concepts: - Reclaim Policy - Access Mode

```
# Task 1. create a volumn for log storage for in a pod
```

```
spec:
```

```
  containers:
```

```
    - name: event-simulator  
      image: kodekloud/event-simulator
```

```
  volumes:
```

```
    - name: log-volume  
      hostPath:  
        path: /var/log/webapp  
        type: Directorymaster $
```

```
# Task 2. create pvc to bound to pv
```

```
# Task 3. use pvc in application pods
```

```
spec:
```

```
  # emit details here
```

```
  volumes:
```

```
    - name: log-volume  
      persistentVolumeClaim:  
        claimName: claim-log-1master $
```

```
# Task 4. try delete pvc? try delete app and then pvc?
```

```
# Persistent Volume

apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0003
spec:
  capacity:
    storage: 5Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Recycle
  storageClassName: slow
  mountOptions:
    - hard
    - nfsvers=4.1
  nfs:
    path: /tmp
    server: 172.17.0.2

# Persistent Volume Claims

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: myclaim
spec:
  accessModes:
    - ReadWriteOnce
  volumeMode: Filesystem
  resources:
    requests:
      storage: 8Gi
  storageClassName: slow
  selector:
    matchLabels:
      release: "stable"
    matchExpressions:
      - {key: environment, operator: In, values: [dev]}
```

## 2.8 Optional topics

---

- static provisioning
- dynamic provisioning
- Stateful sets

## FAQ

---

how to create a cluster?

how to create a pod?

what's the best access pattern

how to use template?

yaml syntax practicing

how to switch context?

what does `READY` column stand for in `get pods` output?

```
# create pod

apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep
spec:
  containers:
  - name: busybox
    image: busybox

# change image name in pod
```

## Reading List

---

[lucassha-CKAD #k8s-tips-service-account CKAD candidate book](#)

## Common Errors

---

### ReplicaSet

The ReplicaSet "replicaset-2" is invalid: spec.template.metadata.labels: Invalid value: map[string]string{"tier":"nginx"}: **selector** does not match template labels

Errors occurring during the provising of pods - ImagePullBackOff - CrashLoopBackOff

## Cmd reference

---

```
kubectl api-versions
kubectl api-resources

# remember shortnames
pv
pvc
ing
...

kk explain pod.spec.nodeSelector
# update resources
kk replace -f example.yaml
kk scale deployment/d1 --replicas=2
kk create --edit

# connect to running container
kubectl -n <namespace> exec -it <pods_name> -- /bin/bash
```